

Comparativa de rendimiento entre Docker y máquina virtual utilizando consultas Cypher en Neo4j

1. Introducción

Este trabajo tiene como objetivo comparar el rendimiento de una base de datos Neo4j ejecutando consultas Cypher en dos entornos distintos: una máquina virtual (VM) con Ubuntu 24.04 y un contenedor Docker ejecutado sobre un host con Ubuntu 22.04. La intención es analizar cuál de los dos entornos ofrece un mejor rendimiento, estabilidad y eficiencia al ejecutar operaciones sobre grafos.

2. Objetivos

- Comparar el rendimiento de Neo4j en una máquina virtual y en un contenedor Docker.
- Medir el tiempo de ejecución de diferentes tipos de consultas Cypher.
- Evaluar la estabilidad de cada entorno a través de la desviación estándar.
- Documentar los resultados obtenidos y extraer conclusiones.

3. ¿Qué es Neo4j y Cypher?

Neo4j es una base de datos orientada a grafos que permite modelar relaciones entre entidades de forma nativa, representando los datos como nodos y relaciones. Cypher es el lenguaje de consulta de Neo4j, similar a SQL pero diseñado para grafos.

4. Entornos de prueba

Máquina virtual

- Sistema operativo: Ubuntu 24.04
- Neo4j instalado manualmente (Community Edition)
- Consultas ejecutadas con `cypher-shell`

Docker

- Host: Ubuntu 22.04
- Contenedor: Imagen oficial `neo4j:latest`
- Neo4j ejecutado con `NEO4J_AUTH=none`

5. Base de datos utilizada

Se simuló una pequeña red social con los siguientes elementos:

Nodos:

- `User`: 50 usuarios con nombre y edad.
- `Post`: 100 publicaciones con contenido y fecha.
- `Tag`: 7 etiquetas temáticas.

Relaciones:

- `(:User)-[:CREATED]->(:Post)`
- `(:Post)-[:HAS_TAG]->(:Tag)`
- `(:User)-[:FOLLOWS]->(:User)`

Los datos se generaron usando consultas Cypher automáticas con `UNWIND` y `rand()` para aleatoriedad controlada.

6. Consultas y metodología de evaluación

Se definieron tres consultas de diferentes niveles de complejidad:

- **Consulta ligera:**

```
MATCH (u:User) RETURN count(u);
```

Cuenta los usuarios. Sirve como referencia para una carga mínima.

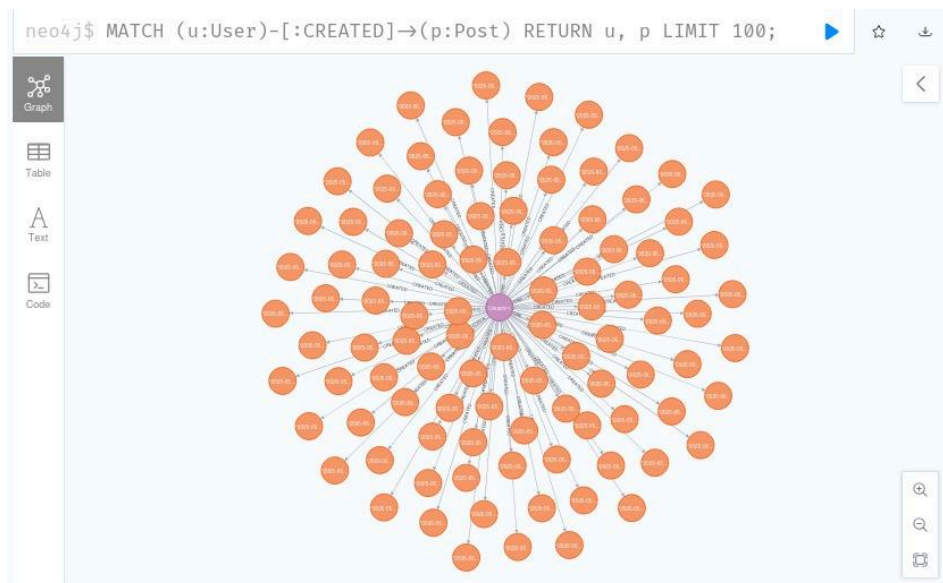
```
neo4j$ MATCH (u:User) RETURN count(u);
```

	count(u)
1	50

- **Consulta intermedia:**

```
MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);
```

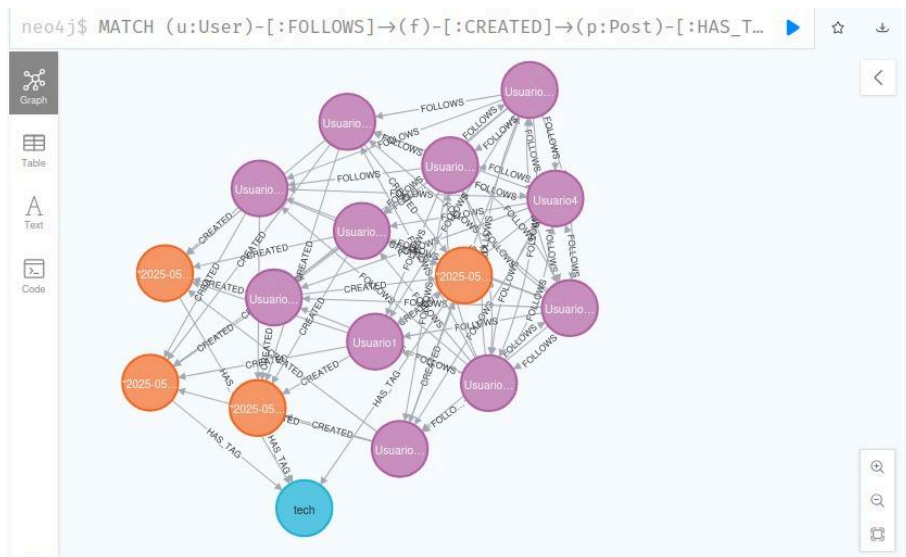
Cuenta las publicaciones creadas por usuarios. Recorre una relación.



- **Consulta compleja:**

```
MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag)
RETURN count(p);
```

Busca publicaciones hechas por usuarios que son seguidos por otros usuarios, y que además están etiquetadas. Luego cuenta esas publicaciones.



- Cada consulta se ejecutó 5 veces seguidas en ambos entornos usando:

time cypher-shell "<consulta>"

Se extrajo el tiempo **real** para cada ejecución. Luego se calculó la media y desviación estándar.

7. Resultados

(Las fotos se encuentran al final del documento)

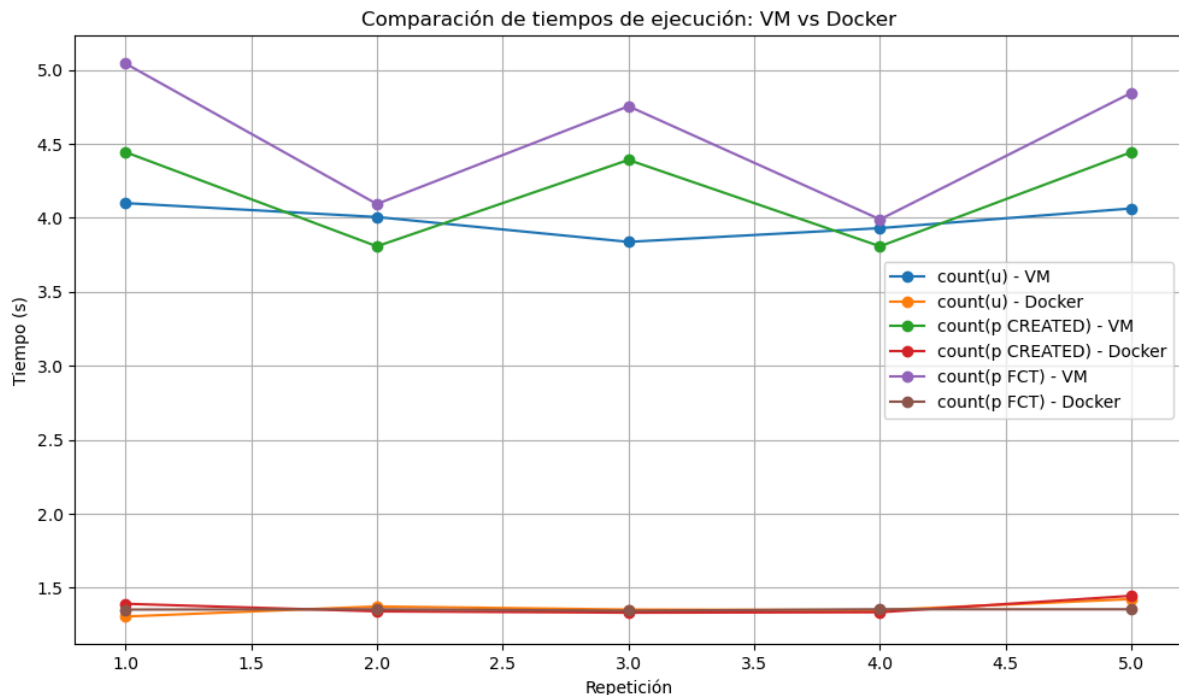
Tabla comparativa:

Consulta	Media VM (s)	Desv. VM	Media Docker (s)	Desv. Docker
count(u)	3.99	0.09	1.35	0.04
count(p) CREATED	4.17	0.21	1.38	0.04
FOLLOWS->CREATED->HAS_ TAG	4.33	0.29	1.38	0.06

8. Análisis comparativo

En las tres consultas, Docker demostró ser significativamente más rápido que la máquina virtual, con mejoras de aproximadamente un 65%. Además, los tiempos fueron más estables en Docker (menor desviación estándar), lo que indica mayor predictibilidad.

Esto puede deberse a que Docker comparte el kernel del sistema operativo y evita la sobrecarga de virtualizar hardware como ocurre en una VM.



9. Conclusión

Docker ofrece una forma más eficiente y rápida de ejecutar Neo4j para este tipo de cargas de trabajo. Es ideal para entornos de desarrollo, pruebas o despliegue continuo. Sin embargo, las máquinas virtuales siguen siendo útiles cuando se requiere aislamiento total del sistema operativo.

10. Repositorio GitHub

El proyecto está disponible en:

https://github.com/beatrizorozco/Docker_vs_VM

Incluye los scripts de configuración, datos generados, consultas ejecutadas y capturas de pantalla de los entornos.

- **VM:**

```
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50       |
+-----+

1 row
ready to start consuming query after 139 ms, results consumed after another 2 ms

real    0m4,100s
user    0m2,650s
sys      0m3,237s
```

```
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50       |
+-----+

1 row
ready to start consuming query after 139 ms, results consumed after another 2 ms

real    0m4,100s
user    0m2,650s
sys      0m3,237s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 600      |
+-----+

1 row
ready to start consuming query after 659 ms, results consumed after another 53 ms

real    0m4,445s
user    0m2,279s
sys      0m3,400s
```

```

bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 6510      |
+-----+

1 row
ready to start consuming query after 450 ms, results consumed after another 126 ms

real    0m5,044s
user    0m2,705s
sys     0m3,562s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50        |
+-----+

1 row
ready to start consuming query after 9 ms, results consumed after another 1 ms

real    0m4,006s
user    0m2,464s
sys     0m3,243s

```

```

bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 600       |
+-----+

1 row
ready to start consuming query after 10 ms, results consumed after another 4 ms

real    0m4,222s
user    0m2,861s
sys     0m3,422s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 6510      |
+-----+

1 row
ready to start consuming query after 2 ms, results consumed after another 26 ms

real    0m4,846s
user    0m2,814s
sys     0m3,296s

```

```

bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50        |
+-----+

1 row
ready to start consuming query after 6 ms, results consumed after another 2 ms

real    0m4,064s
user    0m2,674s
sys     0m3,487s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 600       |
+-----+

1 row
ready to start consuming query after 4 ms, results consumed after another 7 ms

real    0m3,808s
user    0m2,518s
sys     0m2,925s

```

```

bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]
->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 6510      |
+-----+

1 row
ready to start consuming query after 5 ms, results consumed after another 11 ms

real    0m3,909s
user    0m2,605s
sys     0m3,327s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50        |
+-----+

1 row
ready to start consuming query after 6 ms, results consumed after another 3 ms

real    0m3,931s
user    0m2,683s
sys     0m3,126s

```

```

bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 600       |
+-----+

1 row
ready to start consuming query after 12 ms, results consumed after another 5 ms

real    0m3,990s
user    0m2,592s
sys     0m3,302s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]
->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 6510      |
+-----+

1 row
ready to start consuming query after 3 ms, results consumed after another 10 ms

real    0m4,755s
user    0m2,872s
sys     0m3,605s

```

```

bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50        |
+-----+

1 row
ready to start consuming query after 4 ms, results consumed after another 1 ms

real    0m3,838s
user    0m2,702s
sys     0m3,022s
bea@bea-VirtualBox:~$ time cypher-shell -u neo4j -p test "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 600       |
+-----+

1 row
ready to start consuming query after 6 ms, results consumed after another 3 ms

real    0m4,392s
user    0m2,760s
sys     0m3,576s

```


- Docker:

```
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50       |
+-----+

1 row
ready to start consuming query after 36 ms, results consumed after another 1 ms

real    0m1.424s
user    0m0.004s
sys     0m0.010s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 300      |
+-----+

1 row
ready to start consuming query after 75 ms, results consumed after another 6 ms

real    0m1.446s
user    0m0.007s
sys     0m0.012s
```

```
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 1638     |
+-----+

1 row
ready to start consuming query after 88 ms, results consumed after another 8 ms

real    0m1.487s
user    0m0.007s
sys     0m0.012s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50       |
+-----+

1 row
ready to start consuming query after 2 ms, results consumed after another 1 ms

real    0m1.349s
user    0m0.006s
sys     0m0.009s
```

```
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 300      |
+-----+

1 row
ready to start consuming query after 3 ms, results consumed after another 2 ms

real    0m1.345s
user    0m0.002s
sys     0m0.011s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 1638     |
+-----+

1 row
ready to start consuming query after 2 ms, results consumed after another 5 ms

real    0m1.355s
user    0m0.004s
sys     0m0.012s
```

```

bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50       |
+-----+

1 row
ready to start consuming query after 1 ms, results consumed after another 0 ms

real    0m1.332s
user    0m0.005s
sys     0m0.008s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 300      |
+-----+

1 row
ready to start consuming query after 2 ms, results consumed after another 1 ms

real    0m1.335s
user    0m0.005s
sys     0m0.006s

```

```

bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 1638     |
+-----+

1 row
ready to start consuming query after 3 ms, results consumed after another 4 ms

real    0m1.353s
user    0m0.010s
sys     0m0.007s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50       |
+-----+

1 row
ready to start consuming query after 1 ms, results consumed after another 0 ms

real    0m1.351s
user    0m0.008s
sys     0m0.007s

```

```

bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 300      |
+-----+

1 row
ready to start consuming query after 1 ms, results consumed after another 0 ms

real    0m1.374s
user    0m0.008s
sys     0m0.010s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 1638     |
+-----+

1 row
ready to start consuming query after 1 ms, results consumed after another 3 ms

real    0m1.341s
user    0m0.003s
sys     0m0.010s

```

```

bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User) RETURN count(u);"
+-----+
| count(u) |
+-----+
| 50        |
+-----+

1 row
ready to start consuming query after 2 ms, results consumed after another 0 ms

real    0m1.386s
user    0m0.003s
sys     0m0.010s
bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:CREATED]->(p:Post) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 300       |
+-----+

1 row
ready to start consuming query after 1 ms, results consumed after another 1 ms

real    0m1.392s
user    0m0.008s
sys     0m0.011s

```

```

bea@UEA-C289:~$ time sudo docker exec -it neo4j-docker cypher-shell "MATCH (u:User)-[:FOLLOWS]->(f)-[:CREATED]->(p:Post)-[:HAS_TAG]->(t:Tag) RETURN count(p);"
+-----+
| count(p) |
+-----+
| 1638      |
+-----+

1 row
ready to start consuming query after 1 ms, results consumed after another 2 ms

real    0m1.353s
user    0m0.003s
sys     0m0.010s

```